

Using pretrained word embeddings

June 5, 2018

```
In [ ]: # pretrained word embeddings
        ### Hossein Yousefi

In [9]: #Downloading the IMDB data from http://mg.bz/0tIo
import os

imdb_dir = '/home/hossein/Downloads/wordembedding/aclImdb'
train_dir = os.path.join(imdb_dir , 'train')

labels = []
texts = []

for label_type in ['neg', 'pos']:
    dir_name = os.path.join(train_dir, label_type)
    for fname in os.listdir(dir_name):
        if fname[-4:] == '.txt':
            f = open(os.path.join(dir_name, fname))
            texts.append(f.read())
            f.close()
            if label_type == 'neg':
                labels.append(0)
            else:
                labels.append(1)

In [15]: #Tokenizing and vectorizing the data
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
import numpy as np

maxlen = 100
training_samples = 200
validation_samples = 10000
max_words = 10000

tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(texts)
sequences = tokenizer.texts_to_sequences(texts)
```

```

word_index = tokenizer.word_index
print('found %s unique tokens.' % len(word_index))
data = pad_sequences(sequences, maxlen=maxlen)
labels=np.asarray(labels)
print('shape of data tensor :', data.shape)
print('shape of label tensor :', labels.shape)

indices = np.arange(data.shape[0])
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]

x_train = data[:training_samples]
y_train = labels[:training_samples]
x_val = data[training_samples: training_samples + validation_samples]
y_val = labels[training_samples: training_samples + validation_samples]

```

found 88582 unique tokens.
shape of data tensor : (25000, 100)
shape of label tensor : (25000,)

```

In [19]: # Download from https://nlp.stanford.edu/projects/glove/
glove_dir = '/home/hossein/Downloads/wordembedding/glove.6B'
embeddings_index = {}
f = open(os.path.join(glove_dir, 'glove.6B.100d.txt'))
for line in f:
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs
f.closed

```

```

In [20]: embedding_dim = 100

```

```

embedding_matrix = np.zeros((max_words, embedding_dim))
for word, i in word_index.items():
    if i < max_words:
        embedding_vector = embeddings_index.get(word)
        if embedding_vector is not None:
            embedding_matrix[i] = embedding_vector

```

```

In [22]: from keras.models import Sequential
from keras.layers import Embedding, Flatten, Dense

model = Sequential()
model.add(Embedding(max_words, embedding_dim, input_length=maxlen))
model.add(Flatten())

```

```

model.add(Dense(32, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.summary()

model.layers[0].set_weights([embedding_matrix])
model.layers[0].trainable = False

```

```

-----
Layer (type)                Output Shape              Param #
=====
embedding_2 (Embedding)      (None, 100, 100)         1000000
-----
flatten_2 (Flatten)          (None, 10000)             0
-----
dense_3 (Dense)              (None, 32)                320032
-----
dense_4 (Dense)              (None, 1)                 33
=====
Total params: 1,320,065
Trainable params: 1,320,065
Non-trainable params: 0
-----

```

```

In [23]: #training
         model.compile(optimizer='rmsprop',
                       loss='binary_crossentropy',
                       metrics=['acc'])
         history = model.fit(x_train, y_train,
                             epochs=10,
                             batch_size=32,
                             validation_data=(x_val, y_val))
         model.save_weights('pre_trained_glove_model.h5')

```

Train on 200 samples, validate on 10000 samples

Epoch 1/10

200/200 [=====] - 2s 9ms/step - loss: 1.9133 - acc: 0.5150 - val_loss

Epoch 2/10

200/200 [=====] - 1s 6ms/step - loss: 0.6362 - acc: 0.6650 - val_loss

Epoch 3/10

200/200 [=====] - 1s 5ms/step - loss: 0.3214 - acc: 0.8700 - val_loss

Epoch 4/10

200/200 [=====] - 1s 5ms/step - loss: 0.4858 - acc: 0.8500 - val_loss

Epoch 5/10

200/200 [=====] - 1s 5ms/step - loss: 0.1740 - acc: 0.9550 - val_loss

Epoch 6/10

200/200 [=====] - 1s 5ms/step - loss: 0.0705 - acc: 1.0000 - val_loss

Epoch 7/10

```

200/200 [=====] - 1s 5ms/step - loss: 0.2936 - acc: 0.8950 - val_loss
Epoch 8/10
200/200 [=====] - 1s 5ms/step - loss: 0.0483 - acc: 1.0000 - val_loss
Epoch 9/10
200/200 [=====] - 1s 6ms/step - loss: 0.0234 - acc: 1.0000 - val_loss
Epoch 10/10
200/200 [=====] - 1s 5ms/step - loss: 0.0182 - acc: 1.0000 - val_loss

```

```
In [25]: import matplotlib.pyplot as plt
```


```

history_dict = history.history
loss_values = history_dict['loss']
val_loss_values = history_dict['val_loss']

epochs = range(1, len(loss_values) + 1)

plt.plot(epochs, loss_values, 'bo', label='Traning loss')
plt.plot(epochs, val_loss_values, 'b', label='Validation loss')
plt.title('training and validation loss')
plt.xlabel('Epoach')
plt.ylabel('loss')
plt.legend()
plt.show()

```




```

In [26]: plt.clf()
         acc_values = history_dict['acc']
         val_acc_values = history_dict['val_acc']

         plt.plot(epochs, acc_values, 'bo', label='Traning acc')
         plt.plot(epochs, val_acc_values, 'b', label='Validation acc')
         plt.title('Training and validation acc')
         plt.xlabel('Epoach')
         plt.ylabel('Accuracy')
         plt.legend()
         plt.show()

```



```

In [28]: #Evaluation on the test data
         test_dir = os.path.join(imdb_dir, 'test')

         labels = []
         texts = []

         for label_type in ['neg', 'pos']:
             dir_name = os.path.join(test_dir, label_type)
             for fname in sorted(os.listdir(dir_name)):
                 if fname[-4:] == '.txt':
                     f = open(os.path.join(dir_name, fname))
                     texts.append(f.read())

```

```

f.close()
if label_type == 'neg':
    labels.append(0)
else:
    labels.append(1)
sequences = tokenizer.texts_to_sequences(texts)
x_test = pad_sequences(sequences, maxlen=maxlen)
y_test = np.asarray(labels)
model.load_weights('pre_trained_glove_model.h5')
model.evaluate(x_test, y_test)

```

25000/25000 [=====] - 3s 108us/step

Out[28]: [0.8014758405303956, 0.57724]